

calories_burned_analysis

August 27, 2025

[]:

```
'''
About the data:
"This dataset provides a detailed overview of gym members' exercise routines,
↳physical attributes, and fitness metrics. It contains 973
samples of gym data, including key performance indicators such as heart rate,
↳calories burned, and workout duration. Each entry also
includes demographic data and experience levels, allowing for comprehensive
↳analysis of fitness patterns, athlete progression, and
health trends."

From: https://www.kaggle.com/datasets/valakhorasani/gym-members-exercise-dataset

#####

In this analysis we

    1. test the hypothesis that male and female gym members burn the same
↳amount of calories on average and conduct both
    an equal-tailed and one-sided hypothesis test
and
    2. use linear regression to select features within the pipeline and predict
↳the number of calories burned
    using the rest of the data

'''
```

[48]:

```
'''
1) Hypothesis testing
Question: Do male and female gym members burn the same number of calories on
↳average?


Conducting an Equal-tailed Test
H: = ie. Male and female gym members burn the same number of calories on
↳average.
H: ie. Male and female gym members do not burn the same number of
↳calories on average.
'''
```

```
'''

import pandas as pd
df = pd.read_csv("gym_data.csv")
new = df[['Gender', 'Calories_Burned']].dropna()
male_calories = new[new['Gender'] == 'Male']['Calories_Burned']
female_calories = new[new['Gender'] == 'Female']['Calories_Burned']
```

```
[49]: import seaborn as sns
import matplotlib.pyplot as plt

sns.boxplot(x='Gender', y='Calories_Burned', data=df)
plt.title("Calories Burned by Gender")
plt.show()
```



```
[65]: import scipy.stats as stats

t_stat, pvalue = stats.ttest_ind(male_calories, female_calories,
    equal_var=False)
```

```
# pvalue=np.float64(2.0102943133909015e-06)

print('p-value = ', pvalue)
print('male mean cals burned: %.3f' % male_calories.mean())
print('female mean cals burned: %.3f' % female_calories.mean())
```

```
p-value = 2.0102943133909015e-06
male mean cals burned: 944.456
female mean cals burned: 862.249
```

```
[ ]: '''
Since the p-value is less than/equal to 0.001, we can conclude that there is
↳very strong evidence against the hypothesis that male and
female gym members burn the same amount of calories on average based on the
↳observed data. From the sample means, we can see that male
gym members burn more calories than female gym members on average.

'''
```

```
[ ]: '''
Conducting a One-sided Test
H: = ie. Male and female gym members burn the same number of calories on
↳average.
H: > ie. Male gym members burn more calories than female gym members on
↳average.

'''

if t_stat > 0:
    pvalue_onesided = pvalue/2
else:
    pvalue_onesided = 1 - pvalue/2

print('one-sided test p-value = ', pvalue_onesided)
# one-sided test p-value = 1.0051471566954508e-06

'''
Since the one-sided test p-value is less than/equal to 0.001, we can conclude
↳that there is very strong evidence against the hypothesis that
on average, male and female gym members burn the same amount of calories based
↳on the observed data. Furthermore, it seems reasonable to
assume that male gym members burn more calories than female gym members on
↳average.

'''
```

```
[ ]: '''
2) Predictive modeling
Question: Can we predict calories burned per session using member_
↳ characteristics and workout data?

Hypotheses:
- Male gym members burn more calories on average than female gym members
- Of the workout types, Cardio is correlated with burning the most calories
- A longer session duration is correlated with burning more calories
- A higher water intake may be correlated with burning more calories

Generally unsure about Age, BPM, and Workout Frequency as various factors could_
↳ cause competing trends among these features.

'''
```

```
[54]: categorical_features = ['Gender', 'Workout_Type']
numerical_features = ['Age', 'Weight (kg)', 'Height (m)', 'Max_BPM', 'Avg_BPM', _
↳ 'Resting_BPM', 'Session_Duration (hours)', 'Fat_Percentage',
↳ 'Water_Intake (liters)', 'Workout_Frequency (days/
↳ week)', 'Experience_Level', 'BMI']
target_name = 'Calories_Burned'

X = df[numerical_features + categorical_features]
y = df[target_name]
```

```
[68]: from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler

from sklearn.feature_selection import SelectFromModel
from sklearn.svm import LinearSVR

preprocessor = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numerical_features),
        ("cat", OneHotEncoder(drop='first'), categorical_features),
    ]
)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, _
↳ random_state=0)

reg = Pipeline(
    steps=[("preprocessor", preprocessor),
```

```

        ('feature_selection', SelectFromModel(LinearSVR(epsilon=0.01,
↪max_iter=25000))),
        ("regression", LinearRegression())]
)

reg.fit(X_train, y_train)

scores = cross_val_score(reg, X, y, cv=5)
print("scores.mean() %0.2f" % scores.mean())
print("scores.std() %0.2f" % scores.std())

y_pred = reg.predict(X_test)

```

```

scores.mean() 0.87
scores.std() 0.05

```

```

[ ]: '''
      Thus, we have a 5-fold cross-validated  $R^2$  value of  $0.87 \pm 0.05$ , suggesting that
      ↪the model generalizes well.
      '''

```

```

[57]: from sklearn.metrics import root_mean_squared_error

print("R2 score: %.3f" % reg.score(X_test, y_test)) # score method returns  $R^2$ 
↪for regressors by default
print("RMSE: %.3f" % root_mean_squared_error(y_test, y_pred))

```

```

R2 score: 0.868
RMSE: 97.690

```

```

[ ]: '''
      From the output, the model explains 86.8% of the variance in y_test, the set of
      ↪true calories burned among gym members.
      The average magnitude of the model's prediction errors is 97.7 calories from
      ↪the true value.

      Thus, the model is moderately accurate in its predictive performance. Although
      ↪it captures most of the variability in calories burned, an
      average error of around 98 calories makes it somewhat imprecise for reliably
      ↪estimating individual calorie expenditure among gym members.
      '''

```

```

[69]: def get_pipeline_coefficients(fitted_pipeline):
      """
      Returns a DataFrame of selected features and their regression coefficients
      ↪from a fitted pipeline
      with ColumnTransformer, SelectFromModel, and LinearRegression.
      """

```

```

"""
preprocessor = fitted_pipeline.named_steps['preprocessor']
selector = fitted_pipeline.named_steps['feature_selection']
reg = fitted_pipeline.named_steps['regression']

# Get transformed feature names
feature_names = []
for name, transformer, cols in preprocessor.transformers_:
    if name == 'num':
        feature_names.extend(cols)
    else: # categorical
        feature_names.extend(transformer.get_feature_names_out(cols))

# Map selected features to coefficients
selected_features = [feature_names[i] for i in selector.
get_support(indices=True)]
coef_df = pd.DataFrame({
    'Feature': selected_features,
    'Coefficient': reg.coef_
}).sort_values(by='Coefficient', ascending=False).reset_index(drop=True)

return coef_df

coef_df = get_pipeline_coefficients(reg)
print(coef_df)

```

| | Feature | Coefficient |
|---|--------------------------|-------------|
| 0 | Session_Duration (hours) | 247.548701 |
| 1 | Gender_Male | 85.698388 |
| 2 | Workout_Type_Strength | -4.135166 |
| 3 | Workout_Type_HIIT | -5.855987 |
| 4 | Workout_Type_Yoga | -13.960060 |

```
[ ]: '''
The model's coefficients indicate that Session_Duration (hours) is the
↳strongest predictor of calories burned, with each additional
hour of exercise associated with an increase of approximately 248 calories.
↳Male participants are predicted to burn about 86 more
calories than female participants, holding other factors constant. Among
↳workout types, Yoga has the largest negative effect,
reducing predicted calories burned by around 14 calories relative to Cardio,
↳the baseline, while Strength and HIIT sessions have
smaller negative effects. Overall, the results align with expectations: longer
↳sessions and male participants are associated with
higher calorie expenditure, while all workout types burn fewer calories
↳compared to Cardio.
'''
```

```
'''
```

```
[62]: # Beyond evaluating predictive accuracy: statistical significance testing of
      ↪ the selected features from the model
```

```
import statsmodels.api as sm

convert = reg.named_steps['preprocessor']
X_train_transformed = convert.transform(X_train) # a NumPy array
X_train_tra_df = pd.DataFrame(X_train_transformed, columns=convert.
    ↪ get_feature_names_out())
X_train_tra_df = sm.add_constant(X_train_tra_df) # add col of 1s to fit to
    ↪ intercept

y_train_reset = y_train.reset_index(drop=True)

ols_model = sm.OLS(y_train_reset, X_train_tra_df).fit()
print(ols_model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Calories_Burned    R-squared:                0.979
Model:                  OLS               Adj. R-squared:         0.979
Method:                 Least Squares      F-statistic:             2216.
Date:                  Wed, 27 Aug 2025    Prob (F-statistic):       0.00
Time:                  00:05:03           Log-Likelihood:          -3966.2
No. Observations:      778               AIC:                   7966.
Df Residuals:          761               BIC:                   8046.
Df Model:              16
Covariance Type:       nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                        866.8405     3.784    229.096     0.000
859.413    874.268
num__Age                    -41.6357     1.449   -28.733     0.000
-44.480   -38.791
num__Weight (kg)            -21.8668    12.142    -1.801     0.072
-45.704     1.970
num__Height (m)             14.8800     6.811     2.185     0.029
1.509    28.251
num__Max_BPM                -0.1553     1.451    -0.107     0.915
-3.005     2.694
num__Avg_BPM                90.2643     1.439    62.707     0.000
```

| | | | | | |
|------------------------------------|---------|-------------------|----------|---------|-------|
| 87.439 | 93.090 | | | | |
| num__Resting_BPM | | 2.4179 | 1.446 | 1.673 | 0.095 |
| -0.420 | 5.256 | | | | |
| num__Session_Duration (hours) | | 244.5269 | 2.302 | 106.209 | 0.000 |
| 240.007 | 249.047 | | | | |
| num__Fat_Percentage | | -1.8420 | 2.378 | -0.775 | 0.439 |
| -6.510 | 2.826 | | | | |
| num__Water_Intake (liters) | | -1.5120 | 2.239 | -0.675 | 0.500 |
| -5.906 | 2.882 | | | | |
| num__Workout_Frequency (days/week) | | -0.1692 | 2.639 | -0.064 | 0.949 |
| -5.349 | 5.011 | | | | |
| num__Experience_Level | | 0.6357 | 3.342 | 0.190 | 0.849 |
| -5.926 | 7.197 | | | | |
| num__BMI | | 24.1770 | 11.697 | 2.067 | 0.039 |
| 1.214 | 47.140 | | | | |
| cat__Gender_Male | | 82.9710 | 5.181 | 16.015 | 0.000 |
| 72.801 | 93.141 | | | | |
| cat__Workout_Type_HIIT | | -0.9409 | 4.136 | -0.227 | 0.820 |
| -9.061 | 7.179 | | | | |
| cat__Workout_Type_Strength | | -0.6168 | 3.947 | -0.156 | 0.876 |
| -8.366 | 7.132 | | | | |
| cat__Workout_Type_Yoga | | -6.2578 | 4.013 | -1.559 | 0.119 |
| -14.136 | 1.620 | | | | |
| ===== | | | | | |
| Omnibus: | 29.446 | Durbin-Watson: | 2.007 | | |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 46.676 | | |
| Skew: | 0.309 | Prob(JB): | 7.32e-11 | | |
| Kurtosis: | 4.029 | Cond. No. | 23.2 | | |
| ===== | | | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[70]: # Bonferroni correction
k = 5
alpha = 0.05
alpha_adj = alpha / k

def appendcatg(lst):
    L = []
    for col in lst:
        if col.startswith('Session'):
            L.append('num__' + col)
        else:
            L.append('cat__' + col)
    return L
```



```

model_features_bare = coef_df['Feature'].tolist()
model_features = appendcatg(model_features_bare)

L_regular = []
L_bonferroni = []

for col in model_features:
    if ols_model.pvalues[col] < alpha:
        L_regular.append(col)
    if ols_model.pvalues[col] < alpha_adj:
        L_bonferroni.append(col)

print('alpha = 0.05 valid features: ', L_regular)
print('Bonferroni valid features: ', L_bonferroni)

```

```

alpha = 0.05 valid features:  ['num__Session_Duration (hours)',
'cat__Gender_Male']
Bonferroni valid features:  ['num__Session_Duration (hours)',
'cat__Gender_Male']

```

```

[ ]: '''
    Although the predictive model selected five features, statistical testing
    ↪revealed that only two were individually significant predictors of
    calorie expenditure ( $p < 0.05$ , and Bonferroni corrected). This suggests that
    ↪while the model performs adequately for prediction, not all
    selected features show strong independent statistical evidence of association,
    ↪possibly due to multicollinearity or limited sample size.

    '''

```